

Lawtomate: Towards a monolithic, non-iterative prompt for Cambridge Law exams

Motivation:

Cambridge Law exams will be open-book and online for the 2025 series and are expected to continue in this delivery method for the foreseeable.

I believe that:

1. LLMs can do legal exam work faster than traditional test-takers, without requiring the user to actually understand the law and in a manner which can avoid detection as AI-generated content, all whilst maintaining a high standard.

And thus that:

2. Cambridge Law exams are no longer testing exclusively legal reasoning but also now offer an opt-in option to be assessed on prompt-engineering ability instead, should a student persevere at developing these skills (and be willing to break the rules).

Prevailing sentiment amongst Cambridge lawyers denies claim 1. I want to prove that the underlying capability already exists (claim 1), from which the fairness concerns (claim 2) will follow.

My ultimate aim is to demonstrate to Cambridge that this is a genuine and urgent vulnerability. I think the best way to do this is to build an automation system so simple that it democratises these tools at such a low-level, that institutional change must occur.

I am grateful to be financially supported by the Cambridge University AI Research Society and intellectually supported by my college tutor.

Method/Findings etc.:

It is too close to the 2025 series to reasonably expect Cambridge to have enough time to change their delivery format, so I am holding off on providing information until after exams, given that the goal here is to prevent, rather than enable cheating. If you believe you have a good reason to ask for more information, then please email me (link on my homepage). A good reason might be that you are an academic administrator interested in this area. A good reason is not that you are a Cambridge student offering to pay me for details on how to cheat in your exams.

At a basic level, I am comfortable stating that I built out an automated process using API calls and then translated that into a web-UI process that most Cambridge students should easily be able to implement. I am currently working on pushing improvements that happen during that process upstream towards an initial prompt, with the intention of building a prompt that can do more heavy-lifting.

Despite encouragement to publish in Varsity (Cambridge's premier student news venue) from a member of OpenAI's Intelligence and Investigations team, I am planning to approach the law faculty directly instead.

On Education:

This whole situation is representative of an interesting open question about the purpose of education in the face of increasing LLM capabilities. I wrote above that I think institutional change is necessary but purposely have not specified in which direction. Change might entail a return to closed-book approaches, or alternatively it could entail leaning into LLM-assisted workflows and making classes on how to actually leverage these tools part of the curriculum.

The first approach is problematic because the legal world is being transformed by LLMs at present and would risk cocooning law students from reality in a way which might be ultimately unhelpful.

The latter approach is problematic because it risks removing law from Cambridge Law degrees. My latest work on Cambridge essays sit around the 1st class/2.1 boundary (4.0 – 3.6 gpa) and yet I would likely fail if forced to sit a law exam closed-book. It is unclear to me how LLMs could be meaningfully integrated into assessment structures in a manner which preserves their status as tools and prevents them becoming lifelines.

Currently it seems most sensible to teach on LLMs during the year, whilst returning to a closed-book or, at the very least, supervised open-book exam system. This seems to balance for the best of both worlds whilst providing time to think of a better answer.

On Integrity:

Will I be using my method to cheat in my own exams? No – obviously not. That would defeat both the purpose of this project and my entire education. Besides, four of my five exams are closed-book and supervised and I am arranging supervision for the fifth:

